

Key Constraints

Dr Eugenia Ternovska

Simon Fraser University

Keys: superkeys

We need to distinguish tuples within each relation

*The attribute values of a tuple must be such that they must **uniquely identify** the tuple*

No two tuples of a relation are allowed to have the same values for all attributes¹

A **superkey** is a set of one or more attributes that, taken collectively, allow us to identify a tuple uniquely

- ▶ {ID} in table **Instructor** is a superkey
- ▶ {Name} in table **Instructor** is NOT a superkey

¹Some DBMSs relax this requirement

Keys: superkeys and candidate keys

Let R denotes the set of all attributes of relation r and $K \subseteq R$

Formally, K is a superkey of r if

for any tuples $t_1, t_2 \in r$, if $t_1 \neq t_2$ then $t_1.K \neq t_2.K$

(no two distinct tuples have the same values in all attributes of K)

A superkey may contain extraneous (unnecessary) attributes:

- ▶ $\{ID, Name\}$ in table *Instructor* is a superkey

If K is a superkey, then so is any superset of K

Keys: candidate keys are “minimal” superkeys

A **candidate key** is a minimal superkey, i.e., no proper subset of it is a superkey

There could be several candidate keys

- ▶ Suppose that $\{Name, DepName\}$ are sufficient to distinguish among members of the *Instructor* relation
- ▶ $\{ID\}$ in relation (table) *Instructor* is a candidate key
- ▶ $\{Name, DepName\}$ in relation *Instructor* is a candidate key

But:

- ▶ $\{ID, Name\}$ in table *Instructor* is a NOT a candidate key because it is not minimal (includes *ID*)

Keys: primary keys

A **primary key** is a candidate key (i.e., a minimal set of attributes) **chosen by the database designer** as the main way to identify tuples within a relation

It is customary to underline primary keys of a relation schema:

STUDENT(student-id, student-name, major, status)

In database design, primary keys must be chosen with care: their attributes should never change (e.g., SIN), or change extremely rarely (e.g., internal company IDs may change when two companies merge)

Key constraints on a database

A key (primary, candidate or super) is the **property of the entire relation** rather than of the individual tuples

Any two individual tuples in a relation are **prohibited** from having the same values on the key attributes

Therefore, the designation of a key represents a **constraint** on the real-world enterprise being modelled

Keys are probably the **most basic** and **very essential** database constraints

Any update that effects the values of a key, or violate a key integrity constraint, will result in an error state for the DBMS

Key constraints later in the course

We will come back to key constraints when we talk about Functional Dependencies in databases

Foreign-key constraints: Motivation

Consider the attribute `DepName` in `Instructor` relation

Suppose we also have `Department` relation

It would not make sense for a tuple in `Instructor` to have a value for the attribute `DepName` that does not correspond to a department in the `Department` relation

It makes sense to have that

`Instructor.DepName` is a subset of `Department.DepName`

Foreign-key constraints

A **Foreign-key constraint** from attribute A of relation r_1 to the primary key B of relation r_2 states that,

for any database instance, the value of A for each tuple in r_1 must also be the value of B for some tuple in r_2 . In symbols,

$$r_1.A \subseteq r_2.\underline{B}$$

Attribute set A is called **foreign key** from r_1 referencing r_2

- ▶ Attribute A (**DepName**) in relation r_1 (**Instructor**) is a foreign key from r_1 referencing relation r_2 (**Department**)
- ▶ Note: **DepName** must be the primary key in **Department**

As a result of this constraint,

Instructor.DepName is a subset of **Department.DepName**

Foreign-key constraints later in the course

We will come back to foreign-key constraints when we talk about Inclusion Dependencies in databases

Declaring Keys in SQL: Basic constraints

Keywords:

UNIQUE to declare keys

NOT NULL to disallow null values

PRIMARY KEY key + **not NULL**

FOREIGN KEY to reference attributes in other tables

NULL values are, generally, ignored when checking constraints except for **NOT NULL** and **PRIMARY KEY**

Declaring Keys in SQL: Example

```
CREATE TABLE Account (  
  accnum  VARCHAR(12) UNIQUE,  
  branch  VARCHAR(30),  
  custid  VARCHAR(10),  
  balance NUMERIC(14,2)  
);
```

The following insertion gives an error:

```
INSERT INTO Account VALUES  
(1, 'London',      'cust1', 100),  
(1, 'Edinburgh',  'cust3', 200);
```

The following insertion succeeds:

```
INSERT INTO Account VALUES  
(NULL, 'London',      'cust1', 100),  
(NULL, 'Edinburgh',  'cust3', 200);
```

Compound keys

Keys consisting of more than one attribute must be declared using a different syntax

```
CREATE TABLE Movies (  
  m_title      VARCHAR(30),  
  m_director   VARCHAR(30),  
  m_year       SMALLINT,  
  m_genre      VARCHAR(30),  
  UNIQUE (m_title,m_year)  
);
```

This declares the set {m_title,m_year} as a key for Movies

Primary Keys

Essentially **UNIQUE** + **NOT NULL**

```
CREATE TABLE Account (  
  accnum      VARCHAR(12) PRIMARY KEY,  
  branch      VARCHAR(30),  
  custid      VARCHAR(10),  
  balance     NUMERIC(14,2)  
);
```

same as

```
CREATE TABLE Account (  
  accnum      VARCHAR(12) NOT NULL UNIQUE,  
  branch      VARCHAR(30),  
  custid      VARCHAR(10),  
  balance     NUMERIC(14,2)  
);
```

Foreign keys in SQL (1)

```
CREATE TABLE Customer (  
  custid  VARCHAR(10) PRIMARY KEY  
  name    VARCHAR(20),  
  city    VARCHAR(30),  
  address VARCHAR(30)  
);
```

```
CREATE TABLE Account (  
  accnum  VARCHAR(12),  
  branch  VARCHAR(30),  
  custid  VARCHAR(10) REFERENCES Customer(custid),  
  balance NUMERIC(14,2)  
);
```

Every value for attribute `custid` in `Account` must appear among the values of the **primary key** `custid` in `Customer`

Foreign keys in SQL (2)

General syntax (useful for declaring compound foreign keys)

```
CREATE TABLE <table1> (  
  <attr> <type>,  
  ...  
  <attr> <type>,  
  FOREIGN KEY (<list1>)  
  REFERENCES <table2>(<list2>)  
);
```

where

- ▶ `<list1>` and `<list2>` are lists with the **same number** of attributes
- ▶ attributes in `<list1>` are from table `<table1>`
- ▶ attributes in `<list2>` are **unique** in `<table2>`

Referential integrity and database modifications (1)

Deletion can cause problems with foreign keys

Customer	ID	Name	Account	Number	CustID
	cust1	John		123456	cust1
	cust2	Mary		654321	cust2

where Account.CustID is a **foreign key** for Customer.ID

What happens if one deletes (cust1,John) from Customer?

Three approaches are supported in SQL:

1. Reject the deletion operation
2. Propagate it to Account by deleting also (123456,cust1)
3. “Don’t know” approach: keep the tuple in Account, but set CustID value to **NULL**

Referential integrity and database modifications (2)

All three approaches are supported in SQL

```
CREATE TABLE <table1> (  
  <attr> <type>,  
  ...  
  FOREIGN KEY <list1> REFERENCES <table2>(<list2>)  
  <approach>  
)
```

where <approach> can be:

1. Empty: Reject deletions from <table2> causing the FK to be violated (this is the default when <approach> is not specified)
2. **ON DELETE CASCADE**: Propagate the deletion to <name> (tuples in <table1> that violate the FK will be deleted)
3. **ON DELETE SET NULL**: “Don’t know” approach (the values of the attributes in <list1>, for tuples in <name> that violate the FK, are set to **NULL**)

Acknowledgements

[1] Database Systems: The Complete Book, 2nd Edition Hector Garcia-Molina, Jeffrey D. Ullman, Jennifer Widom Prentice Hall, 2009

[2] Database System Concepts, Seventh Edition Avi Silberschatz, Henry F. Korth, S. Sudarshan McGraw-Hill, March 2019 www.db-book.com

Additional references and resources used in preparation of this course are listed on

<https://canvas.sfu.ca/courses/77505/pages/references-and-resources> or mentioned in slides.